

Honey Encryption:

Security Beyond the Brute-force Bound

Ari Juels
Cornell Tech

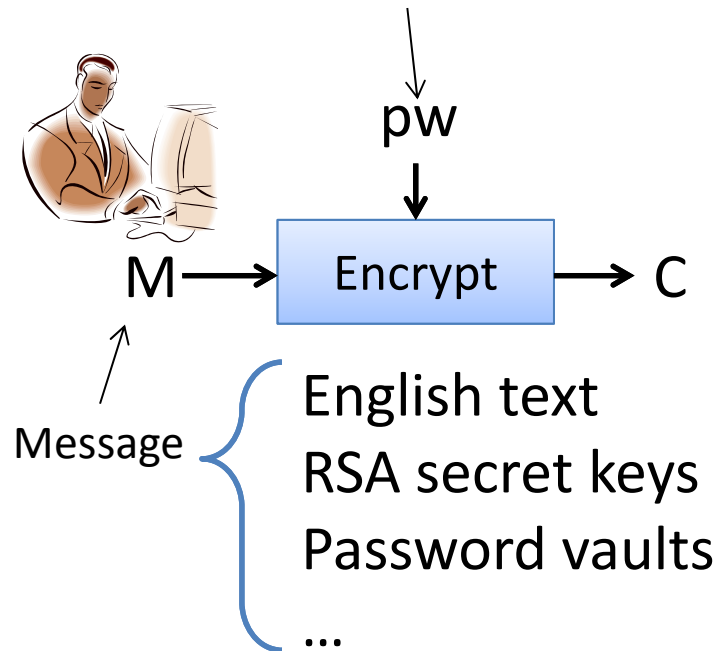
Thomas Ristenpart
University of Wisconsin

Encryption for which decrypting a ciphertext with any number of ***wrong*** keys yields fake, but plausible, plaintexts

Password-based encryption

secret *password* user remembers

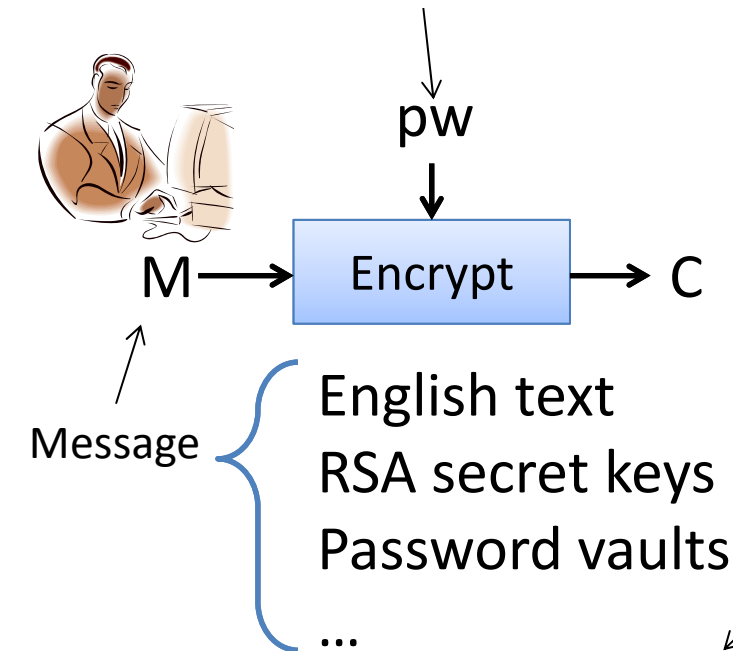
PKCS#5 is dominant standard



Password-based encryption

secret *password* user remembers

PKCS#5 is dominant standard



Encrypt(pw, M)

$\text{salt} \leftarrow \$ \{0,1\}^{128}$

$K \leftarrow H^c(\text{pw} || \text{salt})$

$C \leftarrow K \oplus M$


Return (salt, C)

Decrypt(pw, salt, C)

$K \leftarrow H^c(\text{pw} || \text{salt})$

$M \leftarrow K \oplus C$

Return M



Cryptographic hash function H
(H = SHA-256, SHA-512, etc.)

Common choice is $c = 10,000$

Why hash chains and salts?

Slow down *brute-force attacks*

Internet users ditch “password” password, upgrade to “123456”

Contest for most commonly used terrible password has a new champion.


by Jon Brodtkin - Jan 20 2014, 4:00pm GMT

[Bonneau 2012] studied 69 million Yahoo! Passwords
1.1% of users pick same password

People choose weak passwords

Brute-force attacks

pw likely to fall in short sequence of guesses pw_1, pw_2, pw_3, \dots



Step 1: Trial decrypts

$M_1 \leftarrow \text{Decrypt}(pw_1, \text{salt}, C)$


$M_2 \leftarrow \text{Decrypt}(pw_2, \text{salt}, C)$

$M_3 \leftarrow \text{Decrypt}(pw_3, \text{salt}, C)$

...

Brute-force attacks

pw likely to fall in short sequence of guesses pw_1, pw_2, pw_3, \dots



Say M is unknown **ASCII text** encoded in binary

Many bytes won't be valid ASCII characters, let alone "look" like English text.




Step 1: Trial decryptions

$$M_1 \leftarrow H^c(pw_1 || salt) \oplus C$$

$$M_2 \leftarrow H^c(pw_2 || salt) \oplus C$$

$$M_3 \leftarrow H^c(pw_3 || salt) \oplus C$$

...



Step 2: Find true plaintext

~~$$M_1 = \$\&\%ff1\ 31f^{\wedge}$$~~


~~$$M_2 = hgjk!alc\&ewj$$~~

$$M_3 = copenhagen$$

...

Brute-force attacks

pw likely to fall in short sequence of guesses pw_1, pw_2, pw_3, \dots



Analyses ignore Step 2, conservatively assuming it is trivial for attacker

Say M is unknown **prime number** encoded as integer

- Hash chain slows attack by factor of c
- Salt prevents rainbow tables, provide separation between users

Primality tests will eliminate majority of candidate plaintexts




Step 1: Trial decryptions

$$M_1 \leftarrow H^c(pw_1 || salt) \oplus C$$

$$M_2 \leftarrow H^c(pw_2 || salt) \oplus C$$

$$M_3 \leftarrow H^c(pw_3 || salt) \oplus C$$

...



Step 2: Find true plaintext

~~$$M_1 = 6123410$$~~

~~$$M_2 = 1299827$$~~

~~$$M_3 = 7321162$$~~

...

The Brute-force Bound

Say pw has min-entropy m (most likely password has probability $1/2^m$)

Corollary [BRT12]: **Encrypt** is such that for all IND-CPA adversaries A

$$\frac{t}{c2^m} \leq \text{Adv}(\text{Encrypt}, A) \leq \frac{t}{c2^m}$$


where $t = cq$ for some q is the number of queries to H modeled as a RO, and ignoring small constants and negligible terms

[B12]: most likely password has prob. 1.1% meaning $m \approx 6.5$

So $t > 1,000,000$ makes the above bound close to 1 for $c = 10,000$

- (A) Existing countermeasures help slow down attacks
but only ensure security for high-entropy pw
- (B) Best we can do when targeting IND-CPA

Beyond the brute-force bound?



Key intuition:

Step 2 may be hard for attacker for some message distributions

Say M is uniformly distributed **bit string**

Seems impossible to distinguish!



Step 1: Trial decryptions

$$M_1 \leftarrow H^c(pw_1 || salt) \oplus C$$

$$M_2 \leftarrow H^c(pw_2 || salt) \oplus C$$

$$M_3 \leftarrow H^c(pw_3 || salt) \oplus C$$

...

???

Step 2: Find true plaintext

$$M_1 = 101010101$$

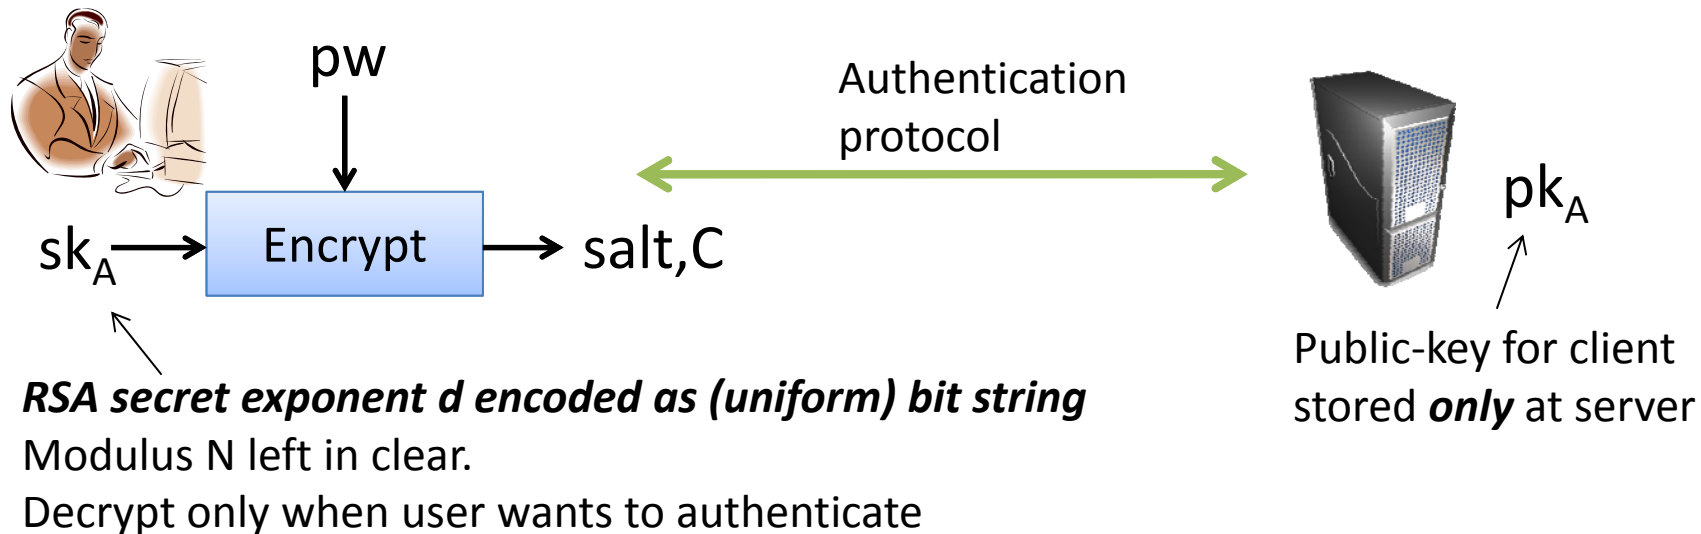
$$M_2 = 100111010$$

$$M_3 = 010101011$$

...

Application: compromise resilience for credentials

[Hoover, Kausik 99]



If attacker just obtains C , best strategy is online attack using M_1, M_2, \dots . Significantly harder to mount than offline attack



Step 1: Trial decryptions

$$M_1 \leftarrow H^c(\text{pw}_1 || \text{salt}) \oplus C$$

$$M_2 \leftarrow H^c(\text{pw}_2 || \text{salt}) \oplus C$$

$$M_3 \leftarrow H^c(\text{pw}_3 || \text{salt}) \oplus C$$

...

Step 2: Find true plaintext

$$M_1 = 101010101$$

$$M_2 = 100111010$$

$$M_3 = 010101011$$


...

???

Decoys in computer security

- In computer security, we have “honey objects”:
 - Honey pots, honeytokens, honey accounts
 - Decoy documents [BHKS09]
 - Kamouflage system [BBBB10]
 - Honeywords for password hashing [JR13]
- Cryptographic camouflage [Hoover, Kausik 99]

We introduce Honey Encryption (HE)



Encryption schemes tailored to specific **message distributions**

Secure in [BRT12] sense (use hash chains and salting)

Provable message-recovery security ***beyond brute-force bound.***

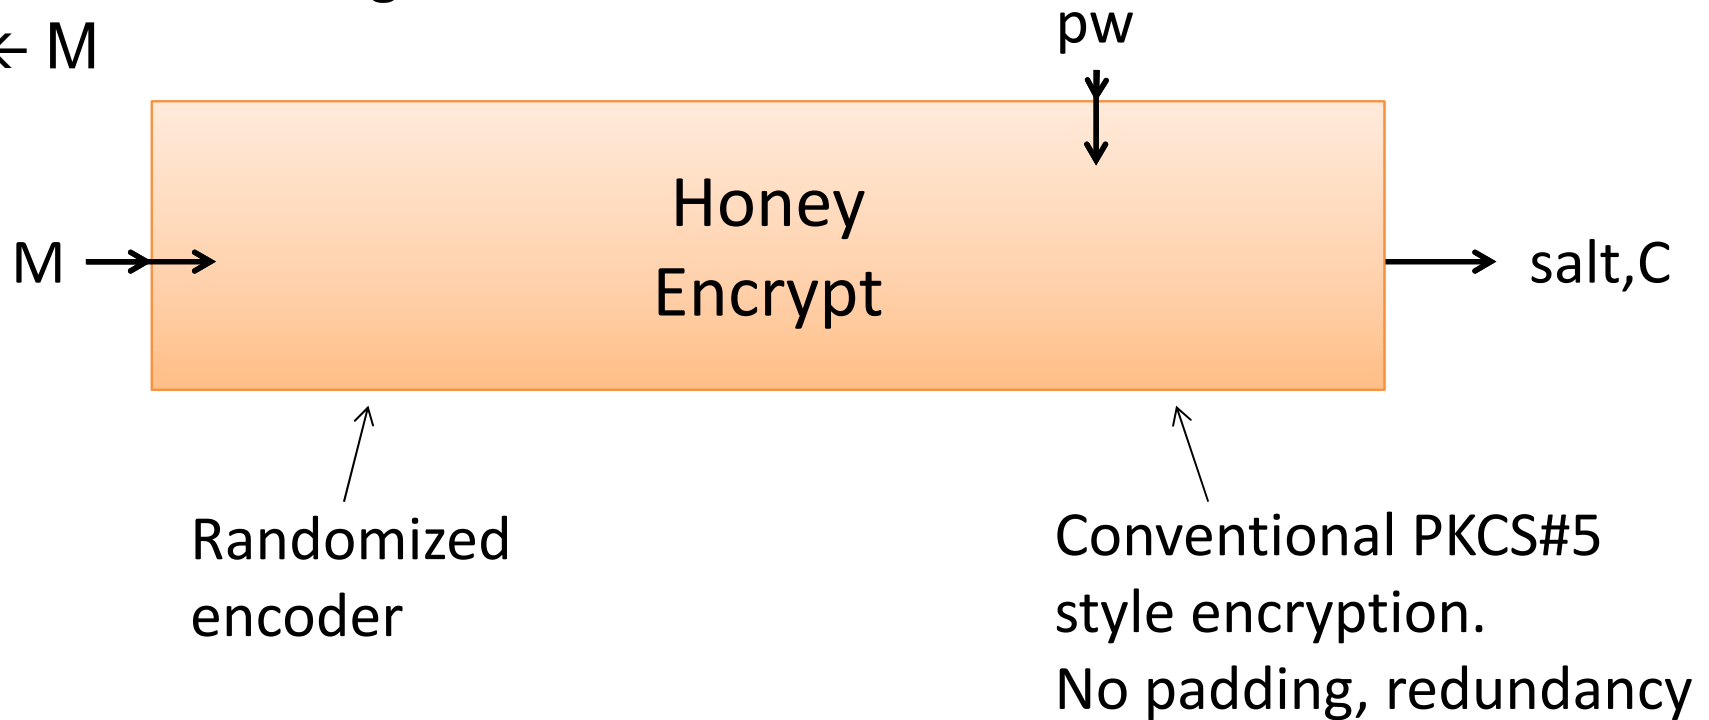
We will show ***optimal security*** in some cases:

$$\Pr[\text{message recovery}] < \frac{1}{2^m}$$

Probability of
guessing
password

A framework for HE schemes


Let M be a message distribution
 $M \leftarrow M$



A framework for HE schemes

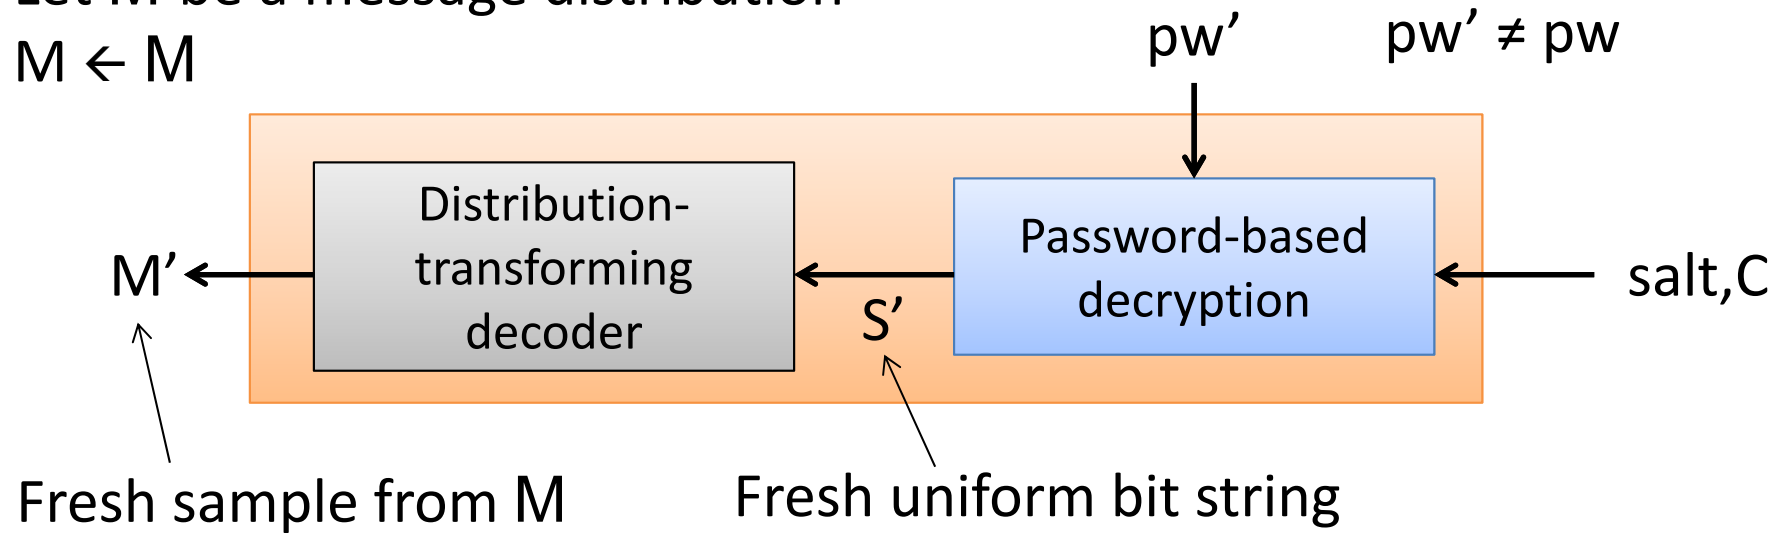
Let M be a message distribution

$M \leftarrow M$



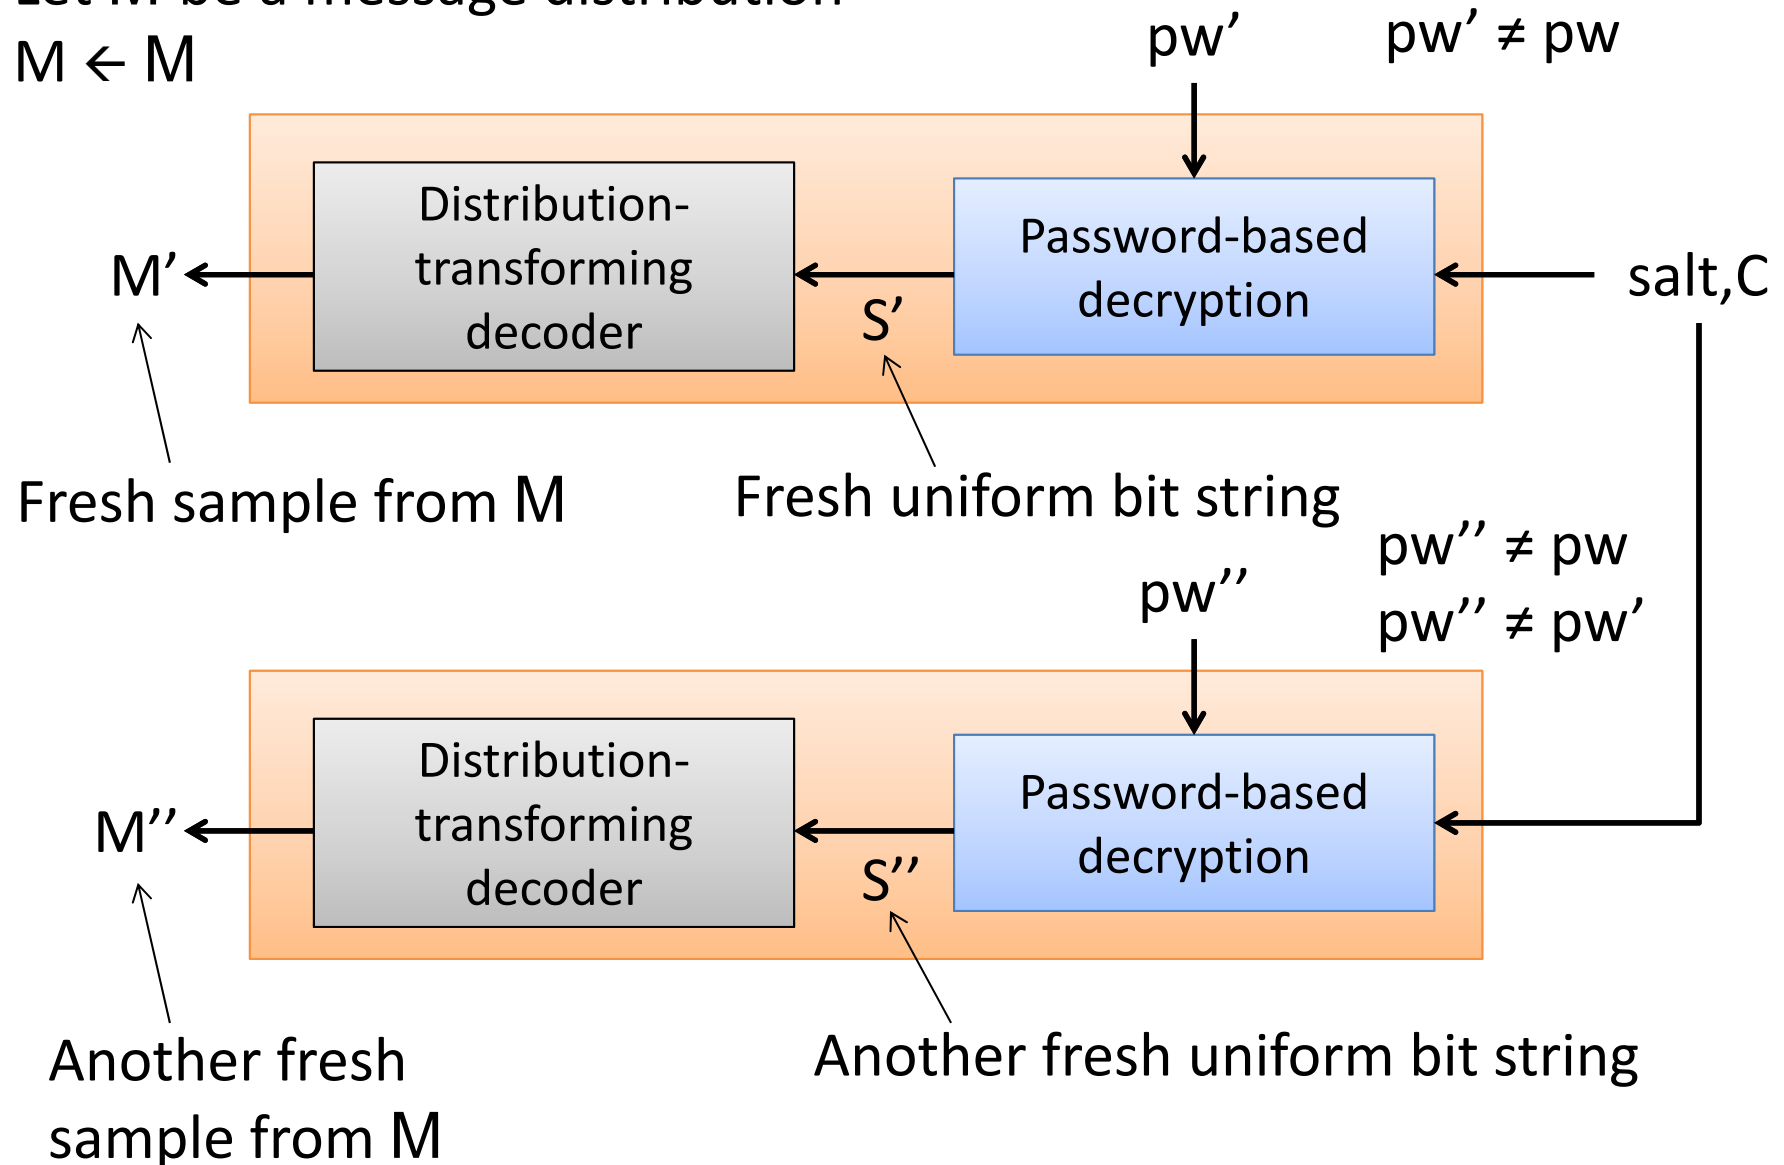
A framework for HE schemes

Let M be a message distribution
 $M \leftarrow M$



A framework for HE schemes


Let M be a message distribution
 $M \leftarrow M$



A framework for HE schemes

Let M be a message distribution

$M \leftarrow M$



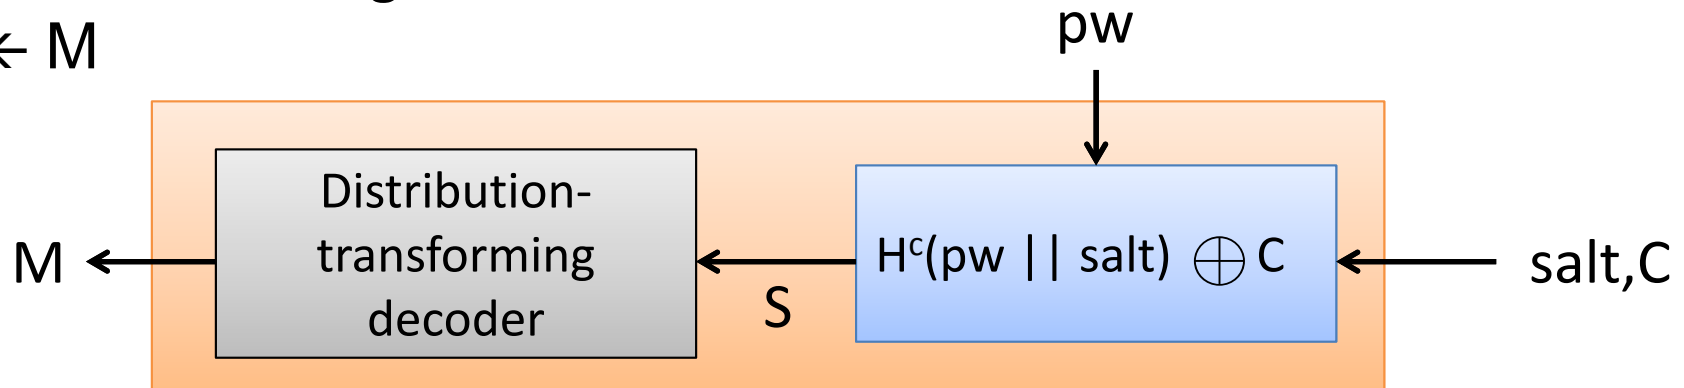
Intuition:

- (1) Decoder is sampler using input as string of randomness
- ✓ (2) Decryption under different keys yields uniform bits

A framework for HE schemes

Let M be a message distribution

$M \leftarrow M$



DTE = (encode, decode) designed for particular M
encode randomized **decode** deterministic

Toy example M

Message	Probability
eurocrypt	1/4
tivoliarden	1/2
Copenhagen	1/4

encode(M)

If $M = \text{tivoliarden}$ then $b \leftarrow \{0,1\}$; Return $0b$

If $M = \text{eurocrypt}$ then Return 11


If $M = \text{Copenhagen}$ then Return 10

decode via look-up table

Huffman coding without compression

A framework for HE schemes

Let M be a message distribution
 $M \leftarrow M$



DTE = (encode, decode) designed for particular M
encode randomized **decode** deterministic

DTE for M being uniform n -bit prime numbers

Encode(M)

$X_1, \dots, X_t \leftarrow \$ (Z_n)^t$
Find 1st i with X_i prime
 $X_i \leftarrow M$
Return $S = X_1, \dots, X_t$


Decode(S)

$X_1, \dots, X_t \leftarrow S$
Find 1st i with X_i prime
 $M \leftarrow X_i$
Return M

Classic
rejection-
sampling prime
generation

A framework for HE schemes

Let M be a message distribution
 $M \leftarrow M$



DTE = (encode, decode) designed for particular M
encode randomized **decode** deterministic

Many DTEs only approximate correct distribution. Secure if:

$M \leftarrow M$
 $S \leftarrow \$ \text{encode}(M)$
Return (M, S)

\approx

$S \leftarrow \$ \{0,1\}^s$
 $M \leftarrow \text{decode}(S)$
Return (M, S)

Honey encryption so far

- Intuition: decryption with wrong password gives plausible plaintext
- Applications in resilience to compromise of encrypted credentials
- Framework:
 - (1) Distribution-transforming encoders (DTEs)
(More examples in paper!)
 - (2) Conventional password-based encryption

Security for honey encryption

Never worse than existing password-based encryption

Inherit provable security in sense of [BRT12]

We analyze **message recovery (MR) security**

MR game:

$M \leftarrow \$ M$

$pw \leftarrow \$ P$

$salt, C \leftarrow \$ HEnc(pw, M)$

$M' \leftarrow \$ A(salt, C)$

Ret $(M=M')$

M is message distribution

P is password distribution

Example: HE for uniform primes

M is uniform n -bit primes

P has min-entropy m

HE scheme as described before

Thm (informal). For any MR attacker A

$\Pr[\text{wins MR game}] < 1/2^m$

(ignoring smaller terms)

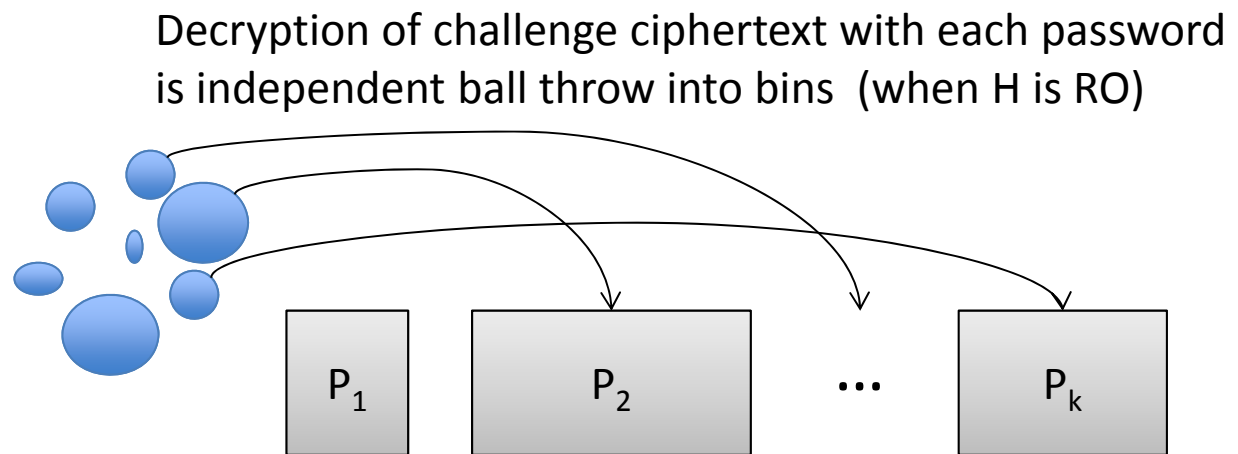
Intuition for proofs

Allow information-theoretic adversaries (also unbounded RO queries)

Adversary outputs most probable message

After applying DTE security, can bound advantage via ***balls-and-bins game***

Balls are passwords
of size equal to their
probability



Adversary's advantage maximized by
picking heaviest bin at end of game

Bins are messages of size
equal to their probability under decode

Expected maximum load $E[L]$ is
expected weight of heaviest bin

Well-studied for some settings

Intuition for proofs

Allow information-theoretic adversaries (also unbounded RO queries)

Adversary outputs most probable message

After applying DTE security, can bound advantage via ***balls-and-bins game***

Balls are passwords
of size equal to their
probability


(Equal weight $1/2^m$ for
uniform distribution)

Adversary's advantage maximized by
picking heaviest bin at end of game

Expected maximum load $E[L]$ is
expected weight of heaviest bin

Well-studied for some settings

Decryption of challenge ciphertext with each password
is independent ball throw into bins (when H is RO)



Bins are messages of size
equal to their probability under decode

(Equal weight $1/2^n$ for uniform distribution)

For prime number HE:

$$k = 2^n \quad \text{and} \quad k^2 \ll 2^m$$

$$\Pr[\text{wins MR game}] < E[L] = 1/2^m + \text{negl}$$

In the paper...

- More DTEs, more HE constructions
- More general balls-and-bins analyses
- Discussion of extensions
 - dealing with password typos
 - detecting online brute-force attacks
- Discussion of limitations of HE

Summary

Def. **Honey Encryption**

Encryption for which decrypting a ciphertext with any number of *wrong* keys yields fake, but plausible, plaintexts

A framework for building and analyzing HE schemes using *Distribution-Transforming Encoders*

Moving forward:

DTEs for more complex distributions

- Password vaults

Further analyses, constructions

- Standard model
- Sharpened balls-and-bins bounds

